

Firewall Report

Akshay Dongaonkar (akd54)

December 5, 2014

1 Objective

I am implementing a Linux based Firewall that mediates access between two interfaces: a pipe to the Internet and a pipe to a protected, internal network. Due to some exigent circumstances of which I've talked to the professor about, most of this design document is not implemented. Instead, I will be talking about my design choices I am thinking of making now; some of which will hopefully be implemented by tomorrow.

2 Project Structure and Pipeline

There are multiple components of this project that I tried to modularize into common library modules for reuse. For example, there is a module called parser that deals with file input and output. It has a function call to read a rules file that codifies firewall rules, it has a function call to parse the `/proc/net/tcp` file to determine if a port is being used or not, and it at one point had code that iterated through a structure that stored log strings and wrote it to a file.

Then, my design decisions diverged. Initially, I wanted to have a multi threaded application that parallelized on the rule. I'd been working on that structure before I had to roll back my code as stated above. I had ingress from an interface working with pcap loops with their own filter and I wanted to extend that to multiple rules. I then wanted to compare this approach versus a stateful firewall which accepted every packet and looked up its validity within a hashtable. If not found, it would compare against the rule set. My multi threaded approach would be very limited in scale with respect to the number of rules, but it may have been easier to maintain and debug.

3 Code and Data Structures

Right now, there is a simple loop and a simple handler that grab packets off of the currently used interface. The handler for the packets captured linearly checks against a filter string. This is not extensible as there is only one rule in effect at all times.

The reason this happened was that I was basing a lot of my code from the libpcap tutorial provided by [Stanford](#). I will be switching to a different architecture if I have enough time.

The current architecture I have in mind is to have multithreaded pcap capture loops that grab packets that conform to their firewall rule. The $n + 1^{th}$ rule will specifically avoid packets captured by the n rule by careful use of the pcap filter expression. So, to make things easier for us, our firewall tries to match to the first rule, going top to bottom.

For a block rule, the captured packet references a special pcap handler that simply does nothing to the packet. For a pass rule, we reference a handler that injects the packet to our internal network interface port. This will give us most of the functionality that we need for our firewall. For our DMZ, we have its rules at the top of our rule set. We then have our capture loop reference a handler that forwards the packet to the DMZ interface port.

However, this will not really scale. We need a way to maintain states of connections so we do not have to handle packet analysis at every packet. For large (elephant) streams, this would be very cumbersome. I am not completely sure as to how I'm going to implement state with my pre-existing code structure.

In addition to handling the flows within the firewall, we still need other structures to implement functionality as required by the project. We need an ARP resolution protocol, we need a NAT function for maintaining duplex connections, and a flexible interface handler.

The NAT function is partially implemented as we parse our system files to determine open TCP ports. We also have functions that guarantee that the port value it returns is not in use at the time of the return. This is done by keeping a bit vector of open ports from the system tcp files. Given this, we can route flows with some NAT mode and store mappings in a hashtable.

The hashtable being used is the same one written for HW2. Considerable work has been done to ensure non-segfaulting behavior within the hashtable so it can be reused in multiple places in this project.

As for the ARP module, where we bind ARP requests for IP's to our address, this is something that can be implemented and unit tested apart from the firewall. This is the next item that will be implemented.

As for the multiple interface handler, this is something that will probably not be implemented. Right now, every interface is hard-coded and switching between them may break dependencies that I do not know of. If it is possible to run, then there will be a configuration file that will multiplex between the different interfaces that will be captured during initialization.

4 Testing Strategy

The testing strategy with my previous code setup was as follows. I was going to be using my HP laptop which runs Ubuntu 14.04 LTS. I was going to capture packets from my `wlan0` interface to feed to my firewall. Once the packets are ingested by the program, the firewall will analyze the packets and match against rules as described in the aforementioned section. My firewall would then route the packet to an output interface.

To verify the rules worked, I would check against packets using Wireshark. A picture of a sample Wireshark capture all is shown in Figure 1.

I would capture the destination interface for my packets. This is very useful as I can see whether my rules accurately work in real time. I can then replicate my test with a dumped

No.	Time	Source	Destination	Protocol	Length	Info
3325	18.925815000	109.201.152.145	192.168.252.2	RTMP	1514	Unknown (0x0)
3326	18.925840000	192.168.252.2	109.201.152.145	TCP	66	37476 > macromedia-fcs [ACK]
3327	18.926237000	109.201.152.145	192.168.252.2	RTMP	1514	Unknown (0x0) Unknown (0x0)
3328	18.926834000	109.201.152.145	192.168.252.2	RTMP	1514	Unknown (0x0)
3329	18.926858000	192.168.252.2	109.201.152.145	TCP	66	37476 > macromedia-fcs [ACK]
3330	18.927219000	109.201.152.145	192.168.252.2	TCP	1514	macromedia-fcs > 37476 [ACK]
3331	18.927767000	109.201.152.145	192.168.252.2	TCP	1514	macromedia-fcs > 37476 [ACK]
3332	18.927797000	192.168.252.2	109.201.152.145	TCP	66	37476 > macromedia-fcs [ACK]
3333	18.928166000	109.201.152.145	192.168.252.2	RTMP	1514	Unknown (0x0) Unknown (0x0)
3334	18.929229000	109.201.152.145	192.168.252.2	RTMP	1514	Unknown (0x0)
3335	18.929253000	192.168.252.2	109.201.152.145	TCP	66	37476 > macromedia-fcs [ACK]
3336	18.934764000	109.201.152.145	192.168.252.2	TCP	1514	macromedia-fcs > 37476 [ACK]
3337	18.935161000	109.201.152.145	192.168.252.2	RTMP	1514	Unknown (0x0)
3338	18.935185000	192.168.252.2	109.201.152.145	TCP	66	37476 > macromedia-fcs [ACK]

▶ Frame 640: 1514 bytes on wire (12112 bits), 1514 bytes captured (12112 bits) on interface 0
 ▶ Ethernet II, Src: Giga-Byt_ee:99:71 (00:0f:ea:ee:99:71), Dst: IntelCor_af:ac:da (8c:a9:82:af:ac:da)
 ▶ Internet Protocol Version 4, Src: 109.201.152.145 (109.201.152.145), Dst: 192.168.252.2 (192.168.252.2)
 ▶ Transmission Control Protocol, Src Port: macromedia-fcs (1935), Dst Port: 37476 (37476), Seq: 583208, Ack: 1, Len: 1448

Figure 1: A sample Wireshark capture on our interface. You can see a streaming TCP connection in the capture.

pcap file that Wireshark generates. The testing then also remains contained to a single machine. This way, I will not have to deal much with external network rules.

However, after looking at the self-study slides and various Piazza posts, I was able to set up the virtual network interface. This will be used to debug the ARP handling module I will try and set up by tomorrow. I will also be using Wireshark to physically check packets arriving and exiting these interfaces to make sure my code works and has easy debugging.

5 Progress

Currently I am behind the amount of work I had done for the progress report before Thanksgiving break. I am still working to get myself up to that stage and get some working implementation of the Firewall ready. I anticipate I will be unable to finish given that there will be less than 24 hours to do an entire semester's worth of project. However, I plan to get as many modular elements of the firewall working so that it is relatively simple to put together.

I have most of my file handling interfaces ready to use. I would like to get the ARP resolution protocol complete, as it is orthogonal in its dependencies to the rest of the firewall and can be shown to work.

Given time, I would like to show multiple rules in work between two set interfaces. That is what I am hoping for in the best case.